Time series models

ARIMA (AutoRegressive Integrated Moving Average)

Pros:

- Handles both trend and seasonality
- Works well for stationary data
- Relatively simple to understand and implement

Cons:

- Assumes linear relationships
- Requires stationary data (may need differencing)
- Limited ability to capture complex patterns

When to use:

Use ARIMA when you have stationary data or data that can be made stationary through differencing. It's good for short to medium-term forecasts where the underlying patterns are relatively stable.

VAR (Vector Autoregression)

Pros:

- Captures interdependencies between multiple time series
- Good for analyzing relationships between variables
- Can handle complex dynamics in multivariate data

Cons:

- Assumes linear relationships
- Can be sensitive to the order of variables
- May overfit with too many parameters

When to use:

Use VAR when you have multiple interrelated time series and want to model their relationships simultaneously. It's particularly useful for economic and financial forecasting where multiple variables influence each other.

SARIMA (Seasonal ARIMA)

Pros:

- Explicitly models seasonality
- Handles both trend and seasonal patterns
- Often performs well on data with clear seasonal components

Cons:

- Can be complex to tune
- Requires stationary data after seasonal differencing
- May struggle with irregular seasonality

When to use:

SARIMA is ideal for time series with strong seasonal patterns, such as monthly sales data or quarterly economic indicators. It's an extension of ARIMA that's specifically designed to handle seasonality.

GARCH (Generalized Autoregressive Conditional Heteroskedasticity)
Pros:

- Models volatility in time series
- Particularly useful for financial data
- Can capture periods of high and low volatility

Cons:

- Focuses on volatility rather than mean prediction
- Can be complex to implement and interpret
- May not perform well in low-volatility environments

When to use:
GARCH is particularly useful for financial time series where volatility clustering is common, such as stock prices or exchange rates. It's often used in combination with other models to forecast both the mean and volatility of a series.

Prophet
Pros:

- Handles missing data and outliers well
- Automatically detects seasonality and holidays
- User-friendly and requires minimal tuning

Cons:

- Can be less accurate than more complex models
- May overfit on smaller datasets
- Limited flexibility for custom components

When to use:
Prophet is great for business forecasting tasks, especially when you have strong seasonal patterns or when dealing with data that has lots of outliers or missing values1.

LSTM (Long Short-Term Memory)
Pros:

- Captures long-term dependencies
- Handles non-linear patterns effectively
- Can work with raw, non-stationary data

Cons:

- Requires large amounts of data
- Computationally intensive
- Can be challenging to interpret

When to use:
LSTM is ideal for complex time series with long-term patterns or when you have a large amount of historical data. It's particularly useful for financial market predictions where patterns can be intricate and non-linear1.

Exponential Smoothing
Pros:

- Simple and intuitive
- Works well with trend and seasonality
- Computationally efficient

Cons:

- May not capture complex patterns
- Can be sensitive to outliers
- Assumes consistent trend and seasonality

When to use:
Exponential smoothing is great for short-term forecasts, especially when you have clear trends or seasonal patterns. It's often used in sales and demand forecasting2.

XGBoost (eXtreme Gradient Boosting)
Pros:

- Handles non-linear relationships well
- Can incorporate external features easily
- Often provides high accuracy

Cons:

- Can overfit if not properly tuned
- Less interpretable than simpler models
- May struggle with long-term dependencies

When to use:
XGBoost is excellent when you have additional features beyond just the time series itself. It's particularly useful in competitions and when you need to squeeze out extra performance.

Choosing the Right Model

- For simple, interpretable forecasts with clear trends or seasonality, start with ARIMA or Exponential Smoothing.
- If you have lots of data and complex patterns, consider LSTM.
- For business forecasting with minimal effort, try Prophet.
- When you have additional features or need high accuracy, look into XGBoost.

References:

- https://www.reddit.com/r/MachineLearning/comments/193672o/d_best_time_series_models_for_forecasting/
- https://www.forecastr.co/blog/time-series-forecasting-model
- https://www.altexsoft.com/blog/time-series-analysis-and-forecasting-novel-business-perspectives/
- https://preset.io/blog/time-series-forecasting-a-complete-guide/
- https://www.influxdata.com/time-series-forecasting-methods/
- https://unit8.com/resources/training-forecasting-models/
- https://www.reddit.com/r/MachineLearning/comments/1cv0hl2/d_foundational_time_series_models_overrated/
- https://www.wallstreetmojo.com/time-series-analysis/